

# Stochastic Gradient Descent Convergence

How do we choose the **learning rate**  $\alpha$  for stochastic gradient descent? Also, how do we **debug stochastic gradient descent** to make sure it is getting as close as possible to the global optimum?

## Checking for convergence

Back when we were using batch gradient descent, our standard way for making sure that gradient descent was converging was we would plot the **optimization cost function** as a function of the **number of iterations**.

- **Batch Gradient Descent:**

- Plot  $J_{train}(\theta)$  as a function of the number of iterations of gradient descent.
- $J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$ 
  - $J_{train}(\theta)$  Should decrease on every iteration
- This works when the training set size was small because we could sum over all examples  $m$ .
  - Doesn't work when we have a massive dataset (e.g.  $m = 300,000,000$  records)

- **With stochastic gradient descent:**

- We don't want to have to pause the algorithm periodically to do a summation over all data
- Moreover, the whole point of stochastic gradient descent is to avoid those whole-data summations

- **Take cost function definition:**

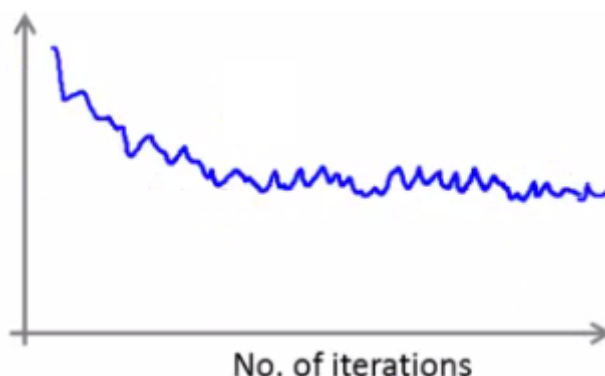
- $cost(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$ 
  - One half the squared error on a single example
- **During learning, compute  $cost(\theta, (x^{(i)}, y^{(i)}))$  before updating  $\theta$  using  $(x^{(i)}, y^{(i)})$ .**
  - i.e. we compute how well the hypothesis is working on the training example
  - Need to do this before we update  $\theta$  because if we did it after  $\theta$  was updated the algorithm would be performing a bit better (because we'd have just used  $(x^{(i)}, y^{(i)})$  to improve  $\theta$ ).

- **Every 1000 iterations (say), plot  $cost(\theta, (x^{(i)}, y^{(i)}))$  averaged over the last 1000 examples processed by algorithm.**

- Gives a running estimate of how well we've done on the last 1000 estimates
- By looking at the plots we should be able to check convergence is happening

Plot  $cost(\theta, (x^{(i)}, y^{(i)}))$ , averaged over the last 1000 (say) examples.

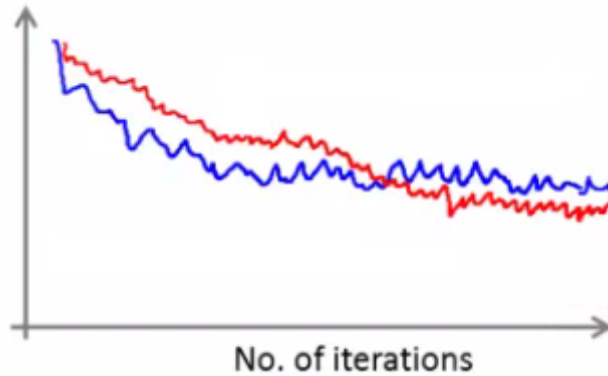
- Might be a bit noisy (1000 examples isn't that much)



If we get the previous figure, that would be a pretty decent run with the algorithm where it looks like the **cost has gone down** and then the plateau that looks kind of flattened out, that is what our cost function looks like then maybe our learning algorithm **has converged**.


- So concretely, the previous figure that's a pretty decent run and the algorithm may have convergence.

If we use a smaller learning rate we may get an even better final solution



With a smaller learning rate, it is possible that we may get a slightly better solution with stochastic gradient descent. That is because stochastic gradient descent **will oscillate and jump around the global minimum**, and it will make smaller random jumps with a **smaller learning rate** (a smaller learning rate means smaller oscillations).

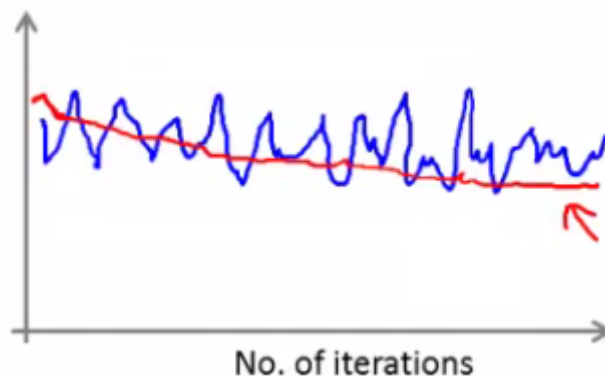
If we average over 1000 examples and 5000 examples we may get a smoother curve



If we increase the number of examples we average over to plot the performance of our algorithm, the plot's line will become smoother.

- This disadvantage of a larger average means we get less frequent feedback

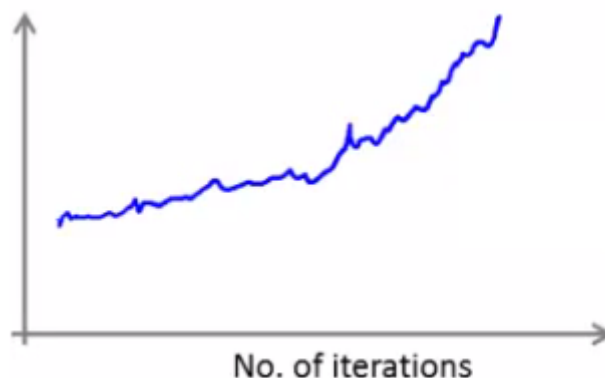
Sometimes we may get a plot that looks like this:



With a very small number of examples for the average, the line will be too noisy and it will be difficult to find the trend.

- So the previous plot looks like cost is not decreasing at all
- But if we then increase to averaging over a larger number of examples we do see this general trend
  - Means the blue line was too noisy, and that noise is ironed out by taking a greater number of entire per average
- So the cost, it may not decrease, even with a large number

If we see a curve that looks like it's increasing then the algorithm may be displaying divergence



- We should use a smaller learning rate


### Stochastic Gradient Descent: Learning rate $\alpha$

In most typical implementations of stochastic gradient descent, **the learning rate  $\alpha$  is typically held constant.** **One strategy for trying to actually converge at the global minimum is to slowly decrease  $\alpha$  over time.**

**For example**  $\alpha = \frac{\text{const1}}{\text{iterationNumber} + \text{const2}}$ . Which means we're guaranteed to converge somewhere.

- We also need to determine const1 and const2

If we tune the parameters well, we can get something like this:



**However, this is not often done because people don't want to have to fiddle with even more parameters.**

**Video Question:** Which of the following statements about stochastic gradient descent are true? Check all that apply.

- Picking a learning rate  $\alpha$  that is very small has no disadvantage and can only speed up learning.

If we reduce the learning rate  $\alpha$  (and run stochastic gradient descent long enough), it's possible that we may find a set of better parameters than with larger  $\alpha$ .

- If we want stochastic gradient descent to converge to a (local) minimum rather than wander or "oscillate" around it, we should slowly increase  $\alpha$  over time.

If we plot  $\text{cost}(\theta, (x^{(i)}, y^{(i)}))$  (averaged over the last 1000 examples) and stochastic gradient descent does not seem to be reducing the cost, one possible problem may be that the learning rate  $\alpha$  is poorly tuned.